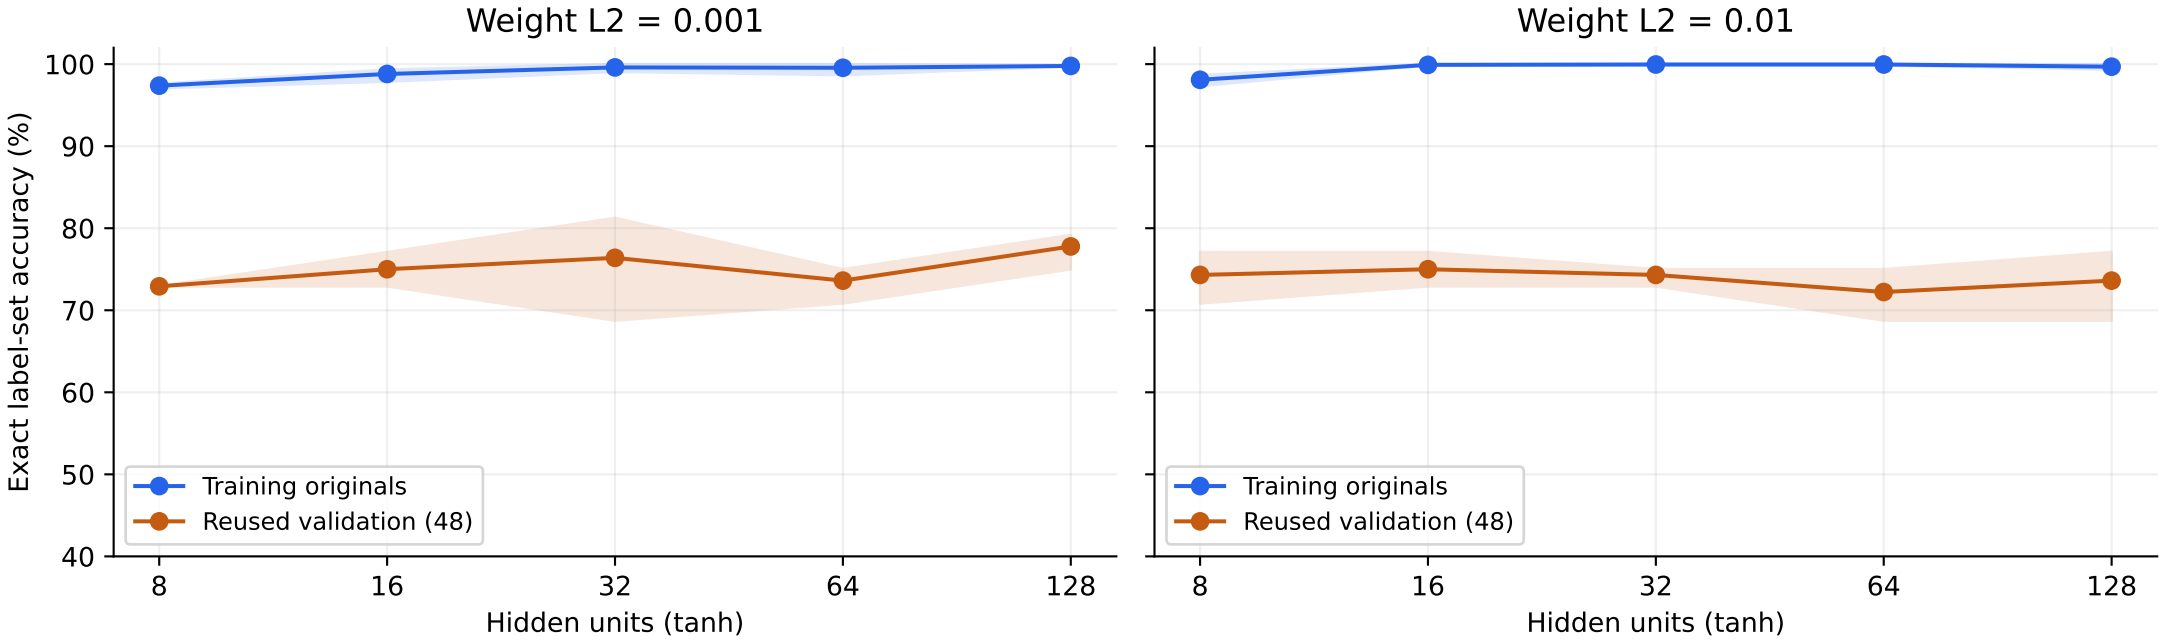
Does increasing capacity close the generalization gap?
Weight L2 = 0.001
Weight L2 = 0.01
Exact label-set accuracy (%)
100
90
80
70
60
50
40
Training originals
Reused validation (48)
8
16
32
64
128
Hidden units (tanh)
Means and min–max bands across three fixed seeds • same optimizer/batch/learning rate • grouped inner stopping